

Event-based Robot Vision

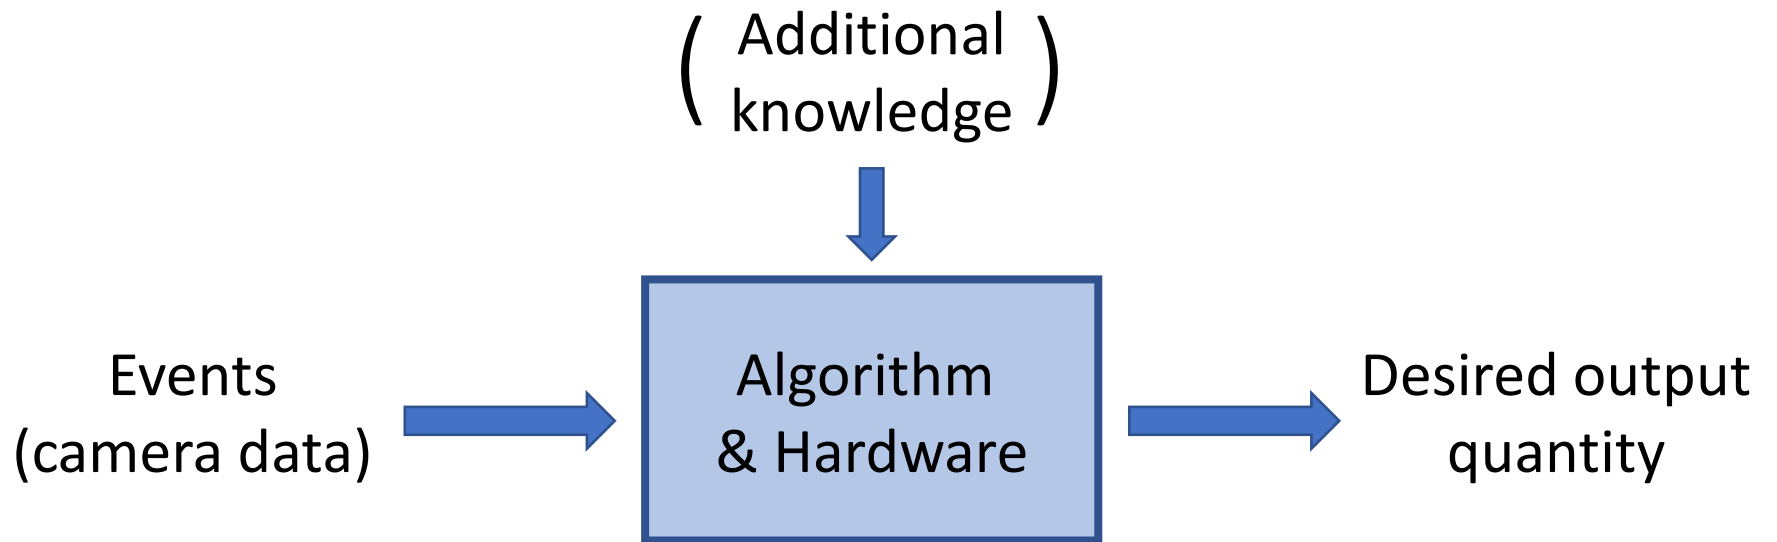
Prof. Dr. Guillermo Gallego
Chair: Robotic Interactive Perception

guillermo.gallego@tu-berlin.de

<http://www.guillermogallego.es>


How to Process the Events?

Overview



The New Paradigm

- Process **one event at a time**
 - How much information does an event carry?
 - The vision system / algorithm needs to have **additional knowledge** (either from external information or built from past events) to assimilate (fuse with) each event.
-
- **Advantages:**
 - **Minimum latency**
 - Asynchronous & sparse
 - **Disadvantages:**
 - Updating the system on a per-event basis may be expensive, depending on the task



The Paradigm: Event-by-event Processing



- **Case Study: Reconstruction of Intensity Image**

Additional knowledge:
Internal “state” $s(\mathbf{x}, t)$
built from past events

Event
 $e_k = (\mathbf{x}_k, t_k, p_k)$


Algorithm
(& Hardware)

Reconstructed
intensity image
 $I(\mathbf{x}, t_k)$



Recall the Event Generation Model

$$\log I(\mathbf{x}, t) - \log I(\mathbf{x}, t - \Delta t) = \pm C$$




Light ($\log I$) has been transduced into asynchronous events...

Given the events, can we recover the absolute intensity?


Let us try to recover the pixel's intensity

- Events represent **intensity changes** \Rightarrow **Integration** should provide absolute intensity



- The recovered signal approximates the original one
- And we cannot see oscillations within the step C (quantization error)
- Additionally, the offset (at $t = 0$) is typically unknown...

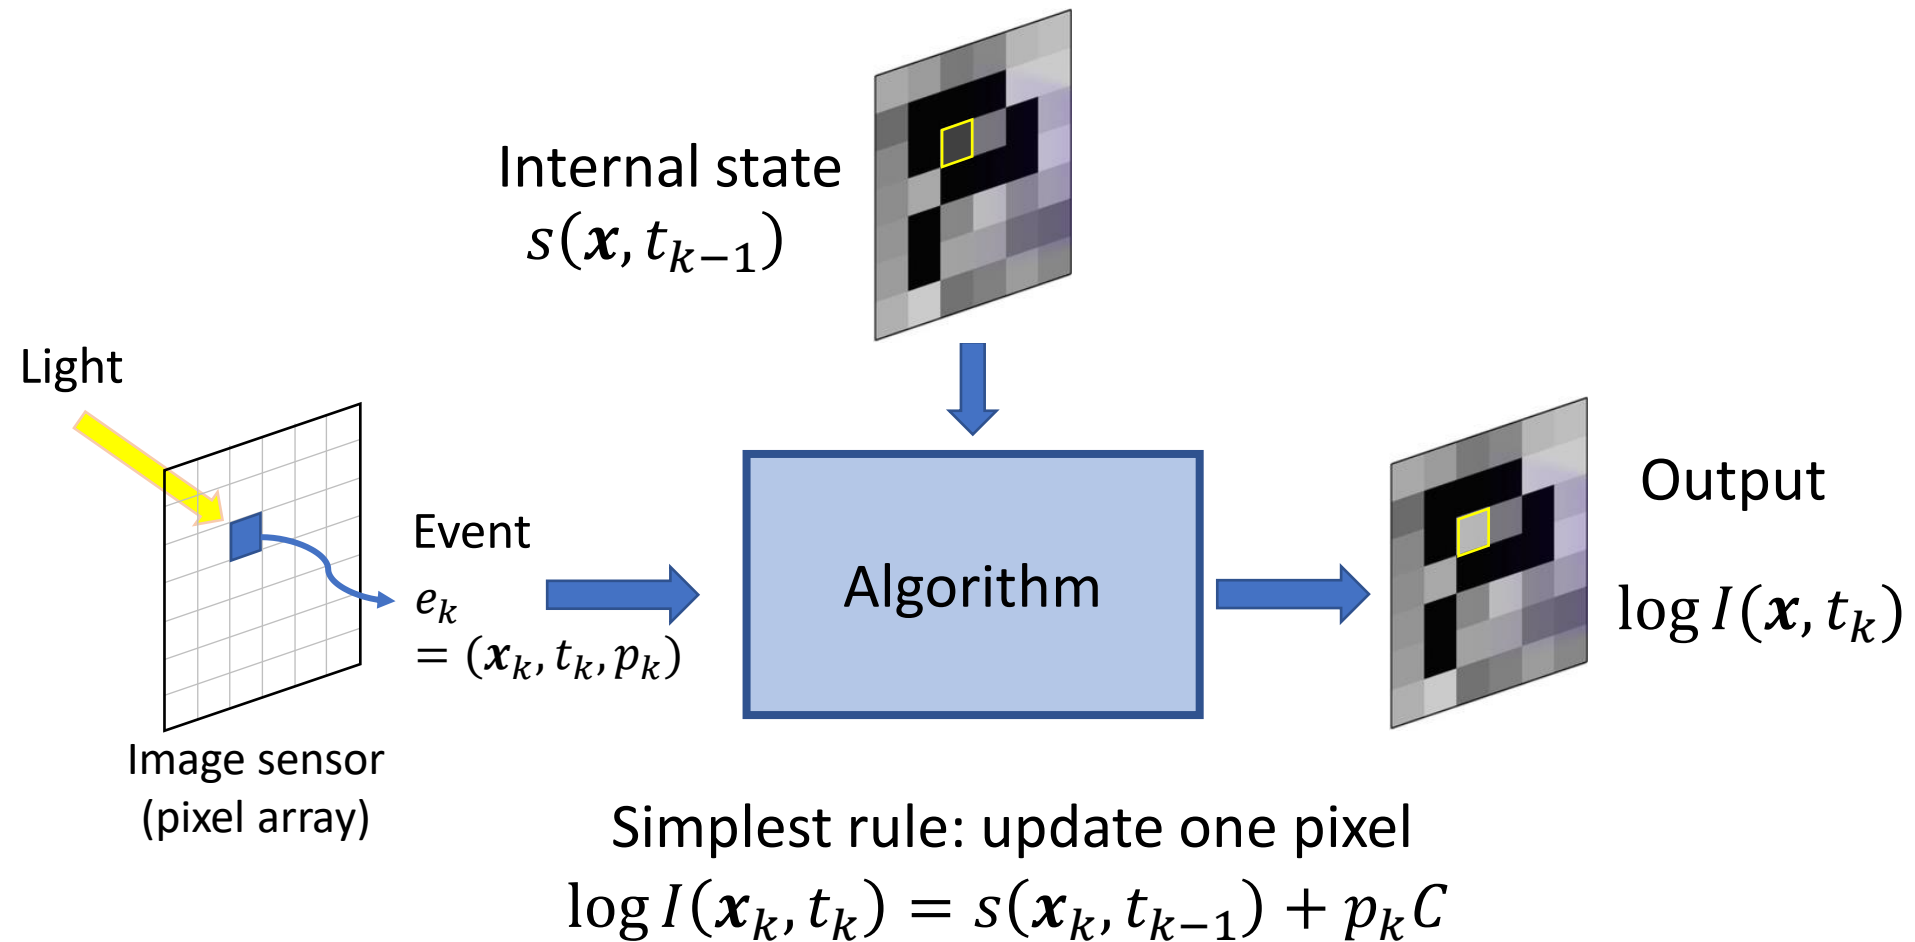
Let us try to recover the pixel's intensity



- Typically the **offset** (at $t = 0$) is unknown...
- That's what happens at one pixel... and we need offsets on all image pixels to make a good (coherent) image

The paradigm: Event-by-event Processing

- **Case Study: Reconstruction of Intensity Image**



Event Integration

$$L(\mathbf{x}, t) = L(\mathbf{x}, 0) + \Delta L(\mathbf{x}; 0, t)$$

Log-intensity
at time t


Log-intensity
at time $t=0$

Increment

intensity in $[0, t]$ is

$$\int_0^t \frac{dL(v)}{dv} dv \approx \sum_k \Delta L_k$$

$$\tilde{L}(\mathbf{x}_k, t_k) = s(\mathbf{x}_k, t_{k-1}) + p_k, \text{ starting from } s=0$$



Direct integration of events,
without starting from $L(\mathbf{x}, 0)$,
cannot recover absolute intensity,
→ Produces incremental intensity ΔL

Notice the **missing offset** $L(\mathbf{x}, 0)$

Event Integration

- Estimated intensity. Intuition:

$$L(t) = L(0) + \underbrace{\int_0^t \frac{dL}{dt}(\tau) d\tau}_{\text{Increment } \Delta L := L(t) - L(0)}$$

Increment $\Delta L := L(t) - L(0)$


$$\log \hat{I}(\mathbf{x}, t) = \log I(\mathbf{x}, 0) + \underbrace{\sum_{0 < t_k \leq t} p_k C \delta(\mathbf{x} - \mathbf{x}_k) \delta(t - t_k)}_{\text{Intensity increment } \Delta \log I}$$

Diagram annotations:


- $\log \hat{I}(\mathbf{x}, t)$: pixel
- $\log I(\mathbf{x}, 0)$: Intensity at $t = 0$ (offset)
- p_k : polarity
- C : Kronecker delta (discrete variable)
- $\delta(\mathbf{x} - \mathbf{x}_k)$: Dirac delta (continuous variable)
- $\delta(t - t_k)$: Dirac delta (continuous variable)

A Replacement for the Integrator?

- The (ordinary) integrator has a very long effect:



- Replace the integrator with a **leaky** one, so that its effect decays over time (i.e., it “forgets”):



Leaky integrator. A first-order filter

A so-called leaky integrator is a first-order filter with feedback. Let's find its transfer function, assuming that the input is $x(t)$ and the output $y(t)$:

$$\frac{dy(t)}{dt} + Ay(t) = x(t)$$

$$\mathcal{L}\left\{\frac{dy(t)}{dt} + Ay(t)\right\} = \mathcal{L}\{x(t)\}$$

where \mathcal{L} denotes application of the [Laplace transform](#). Moving forward:

$$sY(s) + AY(s) = X(s)$$

$$H(s) = \frac{Y(s)}{X(s)} = \frac{1}{s + A}$$

(taking advantage of the Laplace transform's property that $\frac{dy(t)}{dt} \leftrightarrow sY(s)$, assuming that $y(0) = 0$).

This system, with transfer function $H(s)$, has a single pole at $s = -A$. Remember that its frequency response at frequency ω can be found by letting $s = j\omega$:

$$H(j\omega) = \frac{1}{j\omega + A}$$

To get a rough view of this response, first let $\omega \rightarrow 0$:

$$\lim_{\omega \rightarrow 0} H(\omega) = \frac{1}{A}$$

So the system's DC gain is inversely proportional to the feedback factor A . Next, let $\omega \rightarrow \infty$:

$$\lim_{\omega \rightarrow \infty} H(\omega) = 0$$

The system's frequency response therefore goes to zero for high frequencies. **This follows the rough prototype of a lowpass filter.** To answer your other question with respect to its time constant, it's worth checking out the system's time-domain response. Its impulse response can be found by inverse-transforming the transfer function:


$$H(s) = \frac{1}{s + A} \Leftrightarrow e^{-At}u(t) = h(t)$$

where $u(t)$ is the [Heaviside step function](#). This is a very common transform that can often be found in [tables of Laplace transforms](#). This impulse response is an [exponential decay](#) function, which is usually written in the following format:


$$h(t) = e^{-\frac{t}{\tau}}u(t)$$

where τ is defined to be the function's time constant. **So, in your example, the system's time constant is $\tau = \frac{1}{A}$.**

Ordinary integration $\int_0^t f(v)dv \xleftrightarrow{\text{Laplace}} \frac{F(s)}{s}$ has a Fourier response $\frac{1}{s} = \frac{1}{j\omega}$ that blows up at $\omega = 0$.




Leaky integration smooths off the infinite value at $\omega = 0$. Thus the amplitude spectrum $20 \log \frac{1}{|\omega|} = -20 \log |\omega|$, except that it is not ∞ at $\omega = 0$.




$$\alpha = 0.3$$

Leaky integration



$$\tilde{L}(\mathbf{x}_k, t_k) = e^{-\alpha \Delta t_k} s(\mathbf{x}_k, t_{k-1}) + p_k, \text{ starting from } s=0$$



Leaky integration of events can recover \approx absolute intensity, **even without knowing the offset $L(\mathbf{x}, 0)$!**


Isn't it magic?

Temporal leakage “forgets” the initial intensity.


Spatial filtering is **not** needed.

$$\alpha = 2$$

Leaky integration. Short decay



$$\tilde{L}(\mathbf{x}_k, t_k) = e^{-\alpha \Delta t_k} s(\mathbf{x}_k, t_{k-1}) + p_k, \text{ starting from } s=0$$




Too much leakage (forgetting quickly)
Faster decay \rightarrow not integrating well


(even if offset $L(\mathbf{x}, 0)$ is washed out)

$$\alpha = 0.1$$

Leaky integration. Long decay



$$\tilde{L}(\mathbf{x}_k, t_k) = e^{-\alpha \Delta t_k} s(\mathbf{x}_k, t_{k-1}) + p_k, \text{ starting from } s=0$$





Too little leakage (forgetting slowly)
Slower decay \rightarrow integration artefacts

(even if offset $L(\mathbf{x}, 0)$ is washed out)

Noise-Free, Simulated Events

Desktop Scene

Direct (ordinary) Integration



Input: **events** (ON, OFF)


Grayscale frames are just used for visualization.

$$\tilde{L}(\mathbf{x}_k, t_k) = e^{-\alpha \Delta t_k} s(\mathbf{x}_k, t_{k-1}) + p_k$$

starting from $s=0$

$$\alpha = 0$$

Direct (ordinary) Integration



Input: **events** (ON, OFF)



Grayscale frames are just used for visualization.

$$\tilde{L}(\mathbf{x}_k, t_k) = e^{-\alpha \Delta t_k} s(\mathbf{x}_k, t_{k-1}) + p_k$$

starting from $s=0$

$$\alpha = 0.1$$

Leaky integration. Long decay




Input: **events** (ON, OFF)
Grayscale frames are just used for visualization.

$$\tilde{L}(\mathbf{x}_k, t_k) = e^{-\alpha \Delta t_k} s(\mathbf{x}_k, t_{k-1}) + p_k$$

starting from $s=0$

$$\alpha = 0.3$$

Leaky integration






Input: **events** (ON, OFF)
Grayscale frames are just used for visualization.

$$\tilde{L}(\mathbf{x}_k, t_k) = e^{-\alpha \Delta t_k} s(\mathbf{x}_k, t_{k-1}) + p_k$$

starting from $s=0$

$$\alpha = 2$$

Leaky integration. Short decay



Input: **events** (ON, OFF)
Grayscale frames are just used for visualization.



$$\tilde{L}(\mathbf{x}_k, t_k) = e^{-\alpha \Delta t_k} s(\mathbf{x}_k, t_{k-1}) + p_k$$

starting from $s=0$

Real data (noisy)

Bicycle scene

Direct (ordinary) Integration





Input: **events** (ON, OFF)
Grayscale frames are just used for visualization.

$$\tilde{L}(\mathbf{x}_k, t_k) = s(\mathbf{x}_k, t_{k-1}) + p_k$$

starting from $s=0$

$$\alpha = 0.1$$

Leaky Integration. Long decay



Input: **events** (ON, OFF)



Grayscale frames are just used for visualization.

$$\tilde{L}(\mathbf{x}_k, t_k) = e^{-\alpha \Delta t_k} s(\mathbf{x}_k, t_{k-1}) + p_k$$

starting from $s=0$

$$\alpha = 0.3$$

Leaky Integration



Input: **events** (ON, OFF)



Grayscale frames are just used for visualization.

$$\tilde{L}(\mathbf{x}_k, t_k) = e^{-\alpha \Delta t_k} s(\mathbf{x}_k, t_{k-1}) + p_k$$

starting from $s=0$

$$\alpha = 2$$

Leaky Integration. Short decay



Input: **events** (ON, OFF)

Grayscale frames are just used for visualization.

$$\tilde{L}(\mathbf{x}_k, t_k) = e^{-\alpha \Delta t_k} s(\mathbf{x}_k, t_{k-1}) + p_k$$

starting from $s=0$

Complementary Filter

- Combining **frames** and **events** in per-pixel filters.
- **Frames** provide slowly varying information (**low** temporal freq.)
- **Events** provide **high** temporal frequency information

Slowed down 5x



Input: **events** (ON, OFF) and grayscale **frames**

Output of the filter

References

Reading:

- Section 3 of Gallego et al., [Event-based Vision: A Survey](#), TPAMI 2020
- E. Mueggler et al., [The Event-Camera Dataset and Simulator](#), IJRR 2017, page 3.
- Scheerlinck et al., [Continuous-time Intensity Estimation Using Event Cameras](#), ACCV 2018.
- Scheerlinck et al. [Asynchronous Spatial Image Convolutions for Event Cameras](#), RA-L 2019
- Jupyter notebook [demo](#)